Yash Kolekar

Data Scientist | Generative AI, LLMs, RAG, Tool Calling | LangChain/LangGraph, Python, PyTorch

Phone: 8883160707 | Email: yash.kolekar3669@gmail.com | Location: Pune, India

LinkedIn: linkedin.com/in/yash-kolekar-736164216 | GitHub: github.com/yashkolekar1010-cmyk |

Portfolio: yashkolekar1010-cmyk.github.io

Summary

Data Scientist with 2+ years of experience building and deploying generative AI solutions using Large Language Models (LLMs), Retrieval-Augmented Generation (RAG), tool calling and LangChain/LangGraph, with fine-tuning, model evaluation and PyTorch-based machine learning pipelines. Immediate joiner; open to hybrid and remote roles.

Skills

GenAI & LLMs: Generative AI, LLMs, Hugging Face Transformers, Prompt Engineering, Fine-Tuning (LoRA/PEFT)
AI Agents & Retrieval: RAG Pipelines, LangChain/LangGraph, Tool Calling, CrewAI, Embeddings, MCP (familiar)
ML & Deep Learning: Machine Learning (ML), PyTorch, TensorFlow, Scikit-learn, XGBoost, Computer Vision (basic)
Programming: Python, SQL (MySQL, PostgreSQL), ETL, pandas, NumPy, Git, PyTest
Data & Databases: Pinecone, ChromaDB, FAISS, Weaviate, Elasticsearch, MongoDB, Databricks (familiar)
Backend & APIs: API Development, FastAPI, REST APIs, Microservices, OpenAI, Flask & React (basic)
Cloud, DevOps & MLOps: AWS, Azure (familiar), Docker, CI/CD for ML, Jenkins, MLOps, MLflow (basic)
Methods & Practices: Model Evaluation, Model Monitoring, Model Versioning, Performance Metrics, Best Practices

Experience

Consultant – *Giniminds (Client: Infosys)* Sep 2026 – Present

- Implements a generative AI troubleshooting agent for Infosys using Elasticsearch-based search and retrieval.

AI/ML Engineer Apr 2026 – Sep 2026
Webcrypt Technology *Pune, India*

- Implemented generative AI applications with LLMs, including RAG systems, AI agents and chatbots.
- Built RAG pipelines with document chunking, embeddings, semantic search and context optimization.
- Developed dynamic AI agent workflows using LangChain/LangGraph and CrewAI.
- Developed ML pipelines for fine-tuning LLMs with LoRA, and applied prompt engineering to improve outputs.
- Built machine learning models for classification, clustering and forecasting with PyTorch, Scikit-learn, XGBoost.
- Integrated AI models into production via FastAPI, Docker and AWS, with CI/CD and model monitoring.

Implementation Specialist - AI/ML Engineer Jan 2025 – Mar 2026
HotelKey Ind Pvt Ltd *Nagpur, India*

- Developed RAG-based chatbots and NLP applications using Python, LangChain and Hugging Face Transformers.
- Built AI solutions for information retrieval and natural language processing on high-volume enterprise data.
- Deployed FastAPI microservices serving ML models in production through REST APIs.
- Automated build, testing and deployment with Jenkins, Docker, CI/CD pipelines and PyTest.
- Collaborated with technical and non-technical stakeholders on scalable production deployments.

Software Engineer Trainee Feb 2024 – Dec 2024
Starlite Infotech *Pune, India*

- Performed data preprocessing and assisted with Machine Learning experimentation and model evaluation workflows.

Projects

Enterprise RAG + AI Agent Platform *LangChain, LangGraph, OpenAI, Pinecone, FastAPI, AWS*

- Built a RAG pipeline on a Pinecone vector database with LangGraph agents and OpenAI models via FastAPI.

Student Enrollment Assistant Agent *Python, LLM Tool Calling, Agent Loop*

- Designed an agent using tool calling and multi-turn memory, escalating to a human instead of hallucinating.

QuantScan PRO *Python, Scikit-learn, FastAPI, MySQL, React*

- Built a predictive analytics and scoring engine with Scikit-learn models served through FastAPI REST APIs.

Education

Bachelor of Technology (B.Tech) in Computer Science Engineering 2020 – 2024
Solapur University *CGPA: 9.02*

Certifications

AWS Certified Cloud Practitioner (CLF-C02) | GenAI with LangChain & HuggingFace (In Progress) | ML Specialization – Coursera