

Yash Kolekar

☎ 8883160707 ✉ yash.kolekar1010@gmail.com 💻 [yash-kolekar-736164216](https://github.com/yashkolekar1010-cmyk) 🌐 [yashkolekar1010-cmyk](https://github.com/yashkolekar1010-cmyk) 🌐 yashkolekar1010-cmyk.github.io 📍 Pune, India

Summary

Applied AI Engineer with 2+ years delivering production AI automation — integrating LLMs into business workflows via APIs, agentic systems, and backend services. Strong Python background building FastAPI services, RAG pipelines, and LangChain/LangGraph agentic workflows from prototype to production, with hands-on deployment, monitoring, and client-facing delivery experience. Immediate joiner.

Professional Experience

Junior AI Engineer

Apr 2026 – Present

Concept Good Consulting, Pune

- Designed and deployed tailored AI automation solutions end-to-end — translating business requirements into production-grade workflows using Python, FastAPI, and LLM APIs (OpenAI, Claude, Bedrock).
- Built agentic workflows with LangChain, LangGraph, and CrewAI — tool calling, structured outputs, multi-step reasoning — to automate research, reporting, and operational tasks.
- Implemented RAG pipelines using vector databases (Pinecone, ChromaDB, FAISS) for retrieval over enterprise knowledge bases; hardened prototypes into monitored systems on AWS with Docker, Celery, and Redis.
- Produced technical documentation and internal playbooks for deployed solutions to enable reuse; iterated on prompt design and model improvements based on evaluation feedback.

Implementation Specialist — Python / AI Automation

Jan 2025 – Mar 2026

HotelKey Ind Pvt Ltd, Nagpur

- Built and shipped AI automation tools for enterprise clients — RAG retrieval workflows (LangChain, HuggingFace) and GPT-powered chatbots for self-service operations.
- Built FastAPI microservices integrating LLMs into production workflows; improved system interoperability by 80% through reliable API integrations and async processing of 100K+ daily records.
- Automated CI/CD (Jenkins, Docker, Git); raised release stability 35% via PyTest; collaborated with stakeholders to scope requirements, demo solutions, and iterate on feedback.

Projects

Autonomous AI Agent Automation Framework

LangGraph, CrewAI, Python, FastAPI, Redis

- Multi-agent workflow automation system with tool-calling and ReAct reasoning; Celery + Redis async queuing with monitoring dashboard — prototype-to-production delivery.

Enterprise RAG & Retrieval Platform

LangChain, OpenAI, Pinecone, FastAPI, AWS

- Production RAG system with vector search over private knowledge bases; streaming REST APIs, multi-tenant auth, and full deployment documentation on AWS.

QuantScan PRO — AI Automation Scanner

FastAPI, MySQL, Gemini AI, React, Docker

- End-to-end automation pipeline with a 15-filter AI scoring engine over structured market data; automated daily reporting with measurable time savings — owned from prototype through deployment.

Technical Skills

Core: Python (advanced), FastAPI, REST APIs, Async I/O, SQLAlchemy, OOP & Software Design

LLM Platforms & APIs: OpenAI, Claude, Amazon Bedrock, LLaMA, Mistral, Gemini, Prompt Engineering, Structured Outputs, Basic Evaluation Practices

Agentic AI: LangChain, LangGraph, CrewAI, Tool Calling, ReAct, Multi-Agent Systems

RAG & Retrieval: Pinecone, ChromaDB, FAISS, Elasticsearch (familiar), Embeddings, Vector Search, Hybrid Retrieval

Automation & Integration: Workflow Automation, n8n (familiar), API Integrations, Webhooks

Production Engineering: Docker, AWS (Lambda, EC2, S3, Bedrock), Celery, Redis, Async/Queuing Patterns, CI/CD (Jenkins, Git), Monitoring

Client & Delivery: Technical Documentation, Stakeholder Demos, Requirements Discovery, Agile Delivery

Education & Certifications

B.Tech in Computer Science Engineering

2020 – 2024

Solapur University

AWS Cloud Practitioner (CLF-C02) | Generative AI with LangChain & HuggingFace (*In Progress*) | Machine Learning Specialization — Coursera